

# Modular Meta-learning

## Motivation

Previous approaches to meta-learning have focused on finding distributions or initial values of parameters. Our objective is similar, but rather than focusing on transferring information about parameter values, we focus on finding a set of reusable modules that can form components of a solution to a new task, possibly with a small amount of tuning. The authors provide an algorithm, called BounceGrad, which learns a set of modules and then combines them appropriately for a new task.

## Diference between MAML and Modular Meta-learning

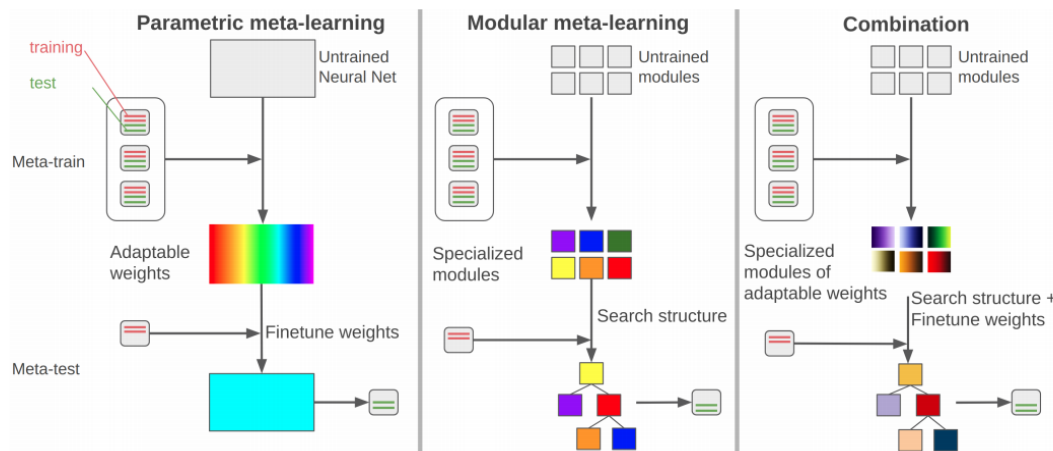
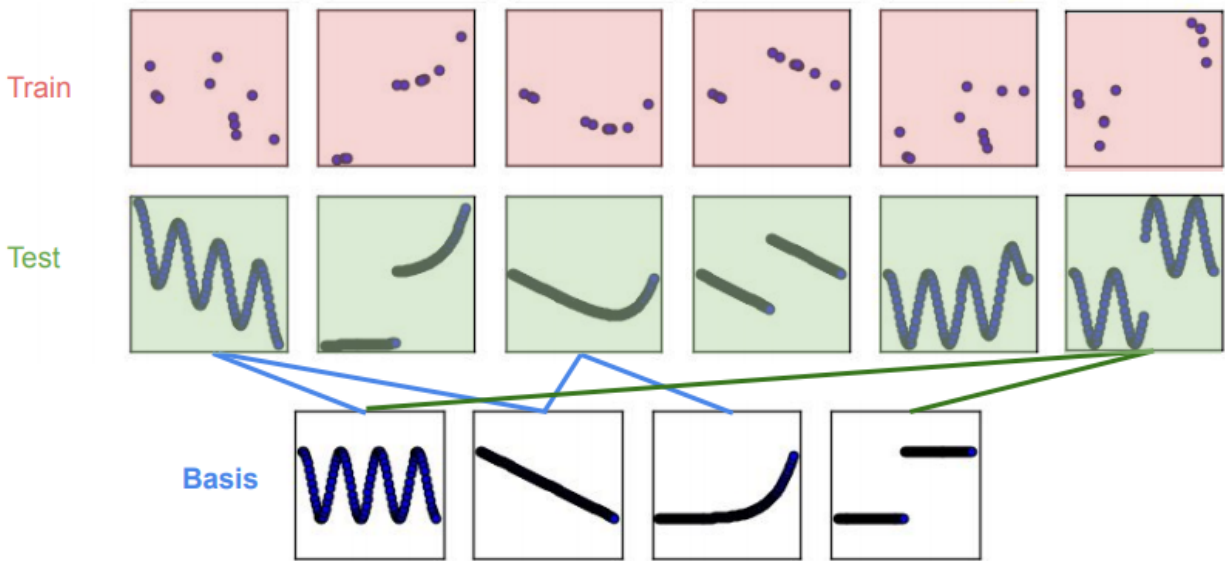


Figure 1 shows that all methods train on a set of related tasks and obtain some flexible intermediate representation. Parametric strategies such as MAML (left) learn a representation that can be quickly adjusted to solve a new task. Our modular meta-learning method (middle) learns a repertoire of modules that can be quickly recombined to solve a new task. A combination of MAML and modular meta-learning (right) learn initial weights for modules that can be combined and adapted for a new task.

## Modular Meta-learning

Here, we give a concrete example to explain what is Modular Meta-learning.



Different from MAML to learn a good initialization, **Modular Meta-learning learns a modular decomposition of characteristics shared by similar tasks.**

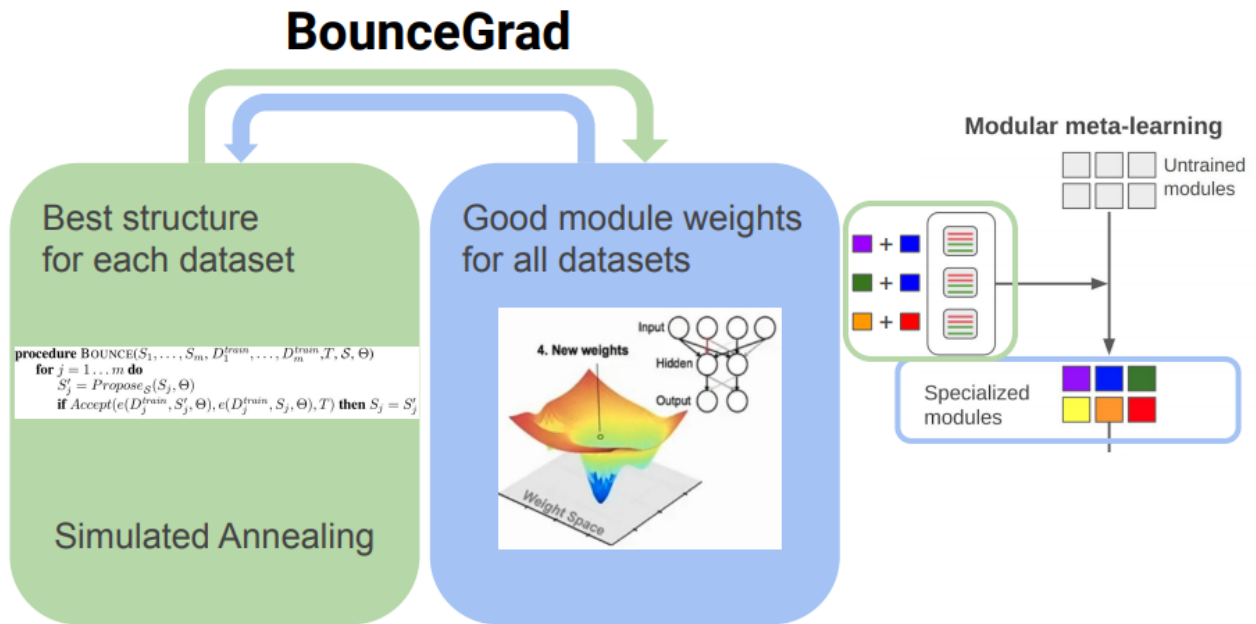
- learn good basis.
- learn which basis to pick and how to compose them together.

## How to learn good modules?

Given the specification of a composition rule and a basis set of modules,  $(C, F, \backslash\theta)$  represents a set of possible functional input-output mappings that will serve as the hypothesis space for the meta-test task.  $F$  is a basis set of modules, which are functions  $f_1, \dots, f_k$ .  $C$  is a compositional scheme for forming complex functions from simpler ones, defined by an initial structure and a set of local modification operations on the structure.

Let  $s$  be the set of possible structures and  $s \text{ in } S$  be a particular structure, generated by  $C$ . The approach has two phases: an off-line meta-learning phase and an on-line meta-test learning phase.

- **Meta-learning phase**, we take training and validation data sets for tasks  $1, \dots, k$  as input and generate a parametrization  $\backslash\theta$  for each module. The objective is to construct modules that will work together as good building blocks for future tasks.
- **Meta-test learning phase**, we take a training data set for the meta-test task as input, as well as  $S$  and  $\backslash\theta$ ; the output is a compositional form  $s \text{ in } S$  which includes a selection of modules  $f_1, \dots, f_m$  to be used in that form. Since  $\backslash\theta$  is already specified, the choice of  $s$  completely determines a mapping from inputs to outputs.



The optimization problems specified previously are in general quite difficult, requiring a mixed continuous-discrete search in a space with many local optima. The authors proposed BounceGrad algorithm (show in above figure), which performs local searches based on a combination of simulated annealing and gradient descent to find approximately optimal solutions to these problems.

This approach considers optimization  $S$  as an inner problem, which has a different aim from existing NAS and meta-learning papers. This approach is more related to dictionary learning. The aim of this approach is to learn pretrained modules ( $\Theta$ ), such that these modules can be composed for new tasks. In this case, the hyperparameters are actually  $\Theta$ , but not  $S$ .